

Decision Trees - Online Shoppers Purchasing Intention

Aim: To build a Decision Tree Classifier model to predict customer behavior based on the features of online shoppers, specifically predicting whether a visitor will make a purchase or not.

Algorithm:

The Decision Tree algorithm selects the best feature to split the data at each node by using criteria like Gini impurity or entropy for classification tasks. It splits the data such that the resulting subsets have the highest possible homogeneity of the target variable.

The splitting criterion for each feature is based on the Gini Index or Information Gain:

Gini Index measures the impurity of a dataset. The Gini impurity of a node is calculated as:

$$Gini(t) = 1 - \sum_{i=1}^k p_i^2$$

Where (p_i) is the proportion of items in class (i) at node (t) .

Entropy measures the disorder or uncertainty in the dataset. It is used in the Information Gain formula, which is computed as:

$$Information\ Gain = Entropy(Parent) - \sum_{i=1}^k \frac{|Subset_i|}{|Parent|} \cdot Entropy(Subset_i)$$

Step 1: Import Libraries

- Import necessary Python libraries that will be used for data manipulation, building the model, and visualizing the results.

Step 2: Load the Dataset

- Load the dataset using pandas `read_csv` function to read the "online_shoppers_intention.csv" file into a DataFrame.

Step 3: Check the Data

- Check the general structure of the dataset using `df.info()` to understand the number of records, data types, and non-null counts.

Step 4: Handle Missing Values

- Check for any missing values in the dataset using `df.isna().sum()` to identify columns with null values.

Step 5: Analyze Target Variable

- Investigate the distribution of the target variable (`Revenue`) to see if it is balanced or imbalanced.

Step 6: Split the Features and Target

- Separate the features (`X`) and the target variable (`y`). The target variable here is `Revenue` , which indicates whether the customer made a purchase (1) or not (0).

Step 7: Encode Categorical Variables

- Use `LabelEncoder` to encode categorical variables (e.g., `Month` and `VisitorType`) into numerical values as Decision Trees require numerical input.

Step 8: Train the Decision Tree Classifier

- Initialize and train a `DecisionTreeClassifier` . Set a maximum depth of 4 to prevent overfitting.

Step 9: Visualize the Decision Tree

- Plot the decision tree to visualize how it makes decisions based on the features. Set appropriate sizes for better visualization.

Step 10: Split Data for Model Evaluation

- Split the dataset into training and testing sets (80% train, 20% test) using `train_test_split` .

Step 11: Fit the Model to Training Data

- Fit the Decision Tree model on the training data.

Step 12: Make Predictions

- Use the trained model to predict the target variable on the test dataset.

Step 13: Evaluate the Model

- Evaluate the model's performance using various classification metrics like accuracy, precision, recall, F1-score, confusion matrix, and classification report.

Import the libraries

```
In [1]: import pandas as pd
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classification_report
import warnings
warnings.filterwarnings("ignore")
```

Load the Dataset

```
In [4]: df = pd.read_csv("../online_shoppers_intention.csv")
```

```
In [5]: df
```

Out[5]:

	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelated	ProductRelated_Duration	BounceRates
0	0	0.0	0	0.0	1	0.000000	0.200000
1	0	0.0	0	0.0	2	64.000000	0.000000
2	0	0.0	0	0.0	1	0.000000	0.200000
3	0	0.0	0	0.0	2	2.666667	0.050000
4	0	0.0	0	0.0	10	627.500000	0.020000
...
12325	3	145.0	0	0.0	53	1783.791667	0.007143
12326	0	0.0	0	0.0	5	465.750000	0.000000
12327	0	0.0	0	0.0	6	184.250000	0.083333
12328	4	75.0	0	0.0	15	346.000000	0.000000
12329	0	0.0	0	0.0	3	21.250000	0.000000

12330 rows × 8 columns

◀

▶

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12330 entries, 0 to 12329
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Administrative         12330 non-null  int64
1   Administrative_Duration 12330 non-null  float64
2   Informational           12330 non-null  int64
3   Informational_Duration  12330 non-null  float64
4   ProductRelated          12330 non-null  int64
5   ProductRelated_Duration 12330 non-null  float64
6   BounceRates             12330 non-null  float64
7   ExitRates              12330 non-null  float64
8   PageValues             12330 non-null  float64
9   SpecialDay             12330 non-null  float64
10  Month                  12330 non-null  object
11  OperatingSystems       12330 non-null  int64
12  Browser                12330 non-null  int64
13  Region                 12330 non-null  int64
14  TrafficType            12330 non-null  int64
15  VisitorType            12330 non-null  object
16  Weekend                12330 non-null  bool
17  Revenue                12330 non-null  bool
dtypes: bool(2), float64(7), int64(7), object(2)
memory usage: 1.5+ MB
```

Check for null values

```
In [7]: df.isna().sum()
```

```
Out[7]: Administrative      0
Administrative_Duration    0
Informational              0
Informational_Duration     0
ProductRelated            0
ProductRelated_Duration   0
BounceRates               0
ExitRates                 0
PageValues                0
SpecialDay                0
Month                     0
OperatingSystems          0
Browser                   0
Region                    0
TrafficType               0
VisitorType               0
Weekend                   0
Revenue                   0
dtype: int64
```

Check target values

```
In [8]: df.Revenue.value_counts()
```

```
Out[8]: Revenue
False    10422
True      1908
Name: count, dtype: int64
```

Split features and target

```
In [9]: y = df[['Revenue']]
y
```

Out[9]:

	Revenue
0	False
1	False
2	False
3	False
4	False
...	...
12325	False
12326	False
12327	False
12328	False
12329	False

12330 rows × 1 columns

```
In [10]: X=df.drop(columns=['Revenue'])
X
```

Out[10]:

	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelated	ProductRelated_Duration	BounceRates
0	0	0.0	0	0.0	1	0.000000	0.200000
1	0	0.0	0	0.0	2	64.000000	0.000000
2	0	0.0	0	0.0	1	0.000000	0.200000
3	0	0.0	0	0.0	2	2.666667	0.050000
4	0	0.0	0	0.0	10	627.500000	0.020000
...
12325	3	145.0	0	0.0	53	1783.791667	0.007143
12326	0	0.0	0	0.0	5	465.750000	0.000000
12327	0	0.0	0	0.0	6	184.250000	0.083333
12328	4	75.0	0	0.0	15	346.000000	0.000000
12329	0	0.0	0	0.0	3	21.250000	0.000000

12330 rows × 17 columns



Encode categorical variables

```
In [11]: enc = LabelEncoder()
X['Month'] = enc.fit_transform(X['Month'])
```

```
X['VisitorType'] = enc.fit_transform(X['VisitorType'])
y['Revenue'] = enc.fit_transform(y['Revenue'])
```

Apply Decision Tree Classifier

```
In [12]: clf = DecisionTreeClassifier(max_depth=4, random_state=42)
```

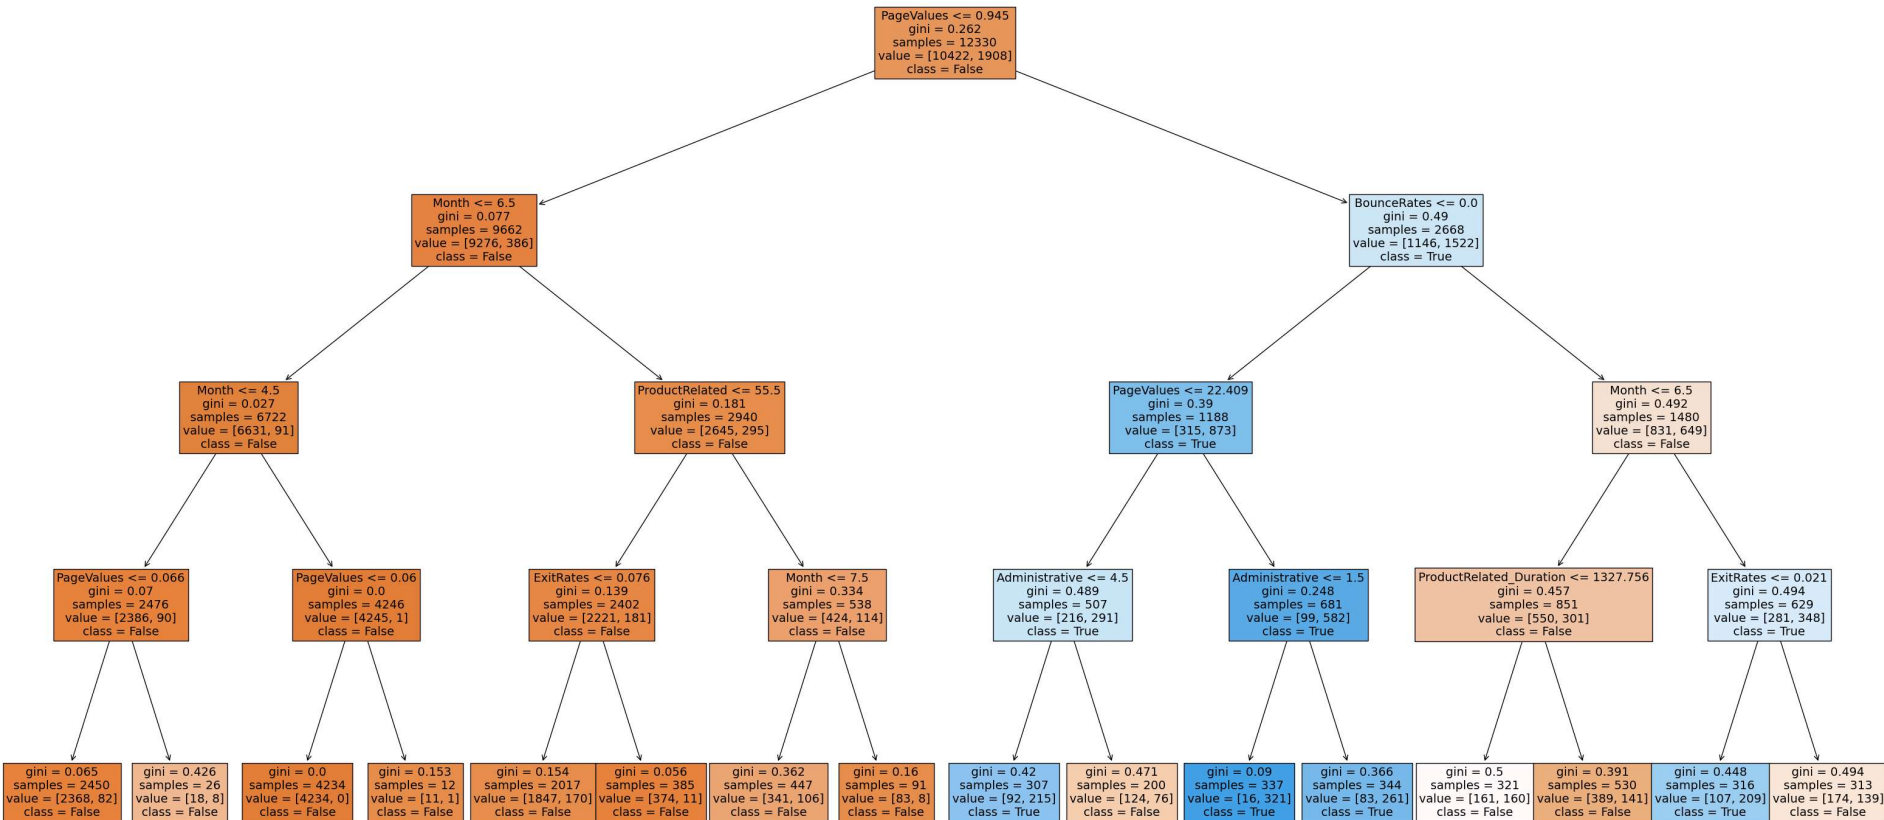
```
In [13]: clf.fit(X,y)
```

```
Out[13]: ▾ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=4, random_state=42)
```

```
In [14]: classNames = enc.classes_.astype(str)
```

Plot the tree

```
In [15]: plt.figure(figsize=(40,20))
plot_tree(clf, feature_names=X.columns, class_names=classNames, filled=True, fontsize=14)
plt.show()
```



Performance metrics

```
In [16]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [17]: clf.fit(X_train, y_train)
```

```
Out[17]: ▾ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=4, random_state=42)
```

```
In [18]: y_pred = clf.predict(X_test)
```

```
In [19]: accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)
```

```
In [20]: print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(class_report)
```

Accuracy: 0.8885
Precision: 0.7036
Recall: 0.5718
F1 Score: 0.6309

Confusion Matrix:
[[1956 99]
[176 235]]

Classification Report:				
	precision	recall	f1-score	support
0	0.92	0.95	0.93	2055
1	0.70	0.57	0.63	411
accuracy			0.89	2466
macro avg	0.81	0.76	0.78	2466
weighted avg	0.88	0.89	0.88	2466

Result

A Decision Tree Classifier was built to predict whether online shop visitors will turn into payinig customer with an accuracy of 88.85%

